

```
In [1]: import collections
import pandas as pd
import numpy as np
import warnings

warnings.filterwarnings(action = 'ignore')

import gensim
from gensim.models import Word2Vec
from sklearn.pipeline import Pipeline
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split, cross_val_score, KFold
from nltk.tokenize import sent_tokenize, word_tokenize
```

```
In [2]: train_raw_df = pd.read_csv('cardataset.csv')
```

```
In [3]: train_raw_df
```

Out[3]:

	Make	Model	Year	Engine Fuel Type	Engine HP	Engine Cylinders	Transmission Type	Driven_Wheels	Number of Doors	
0	BMW	1 Series M	2011	premium unleaded (required)	335.0	6.0	MANUAL	rear wheel drive	2.0	Fac
1	BMW	1 Series	2011	premium unleaded (required)	300.0	6.0	MANUAL	rear wheel drive	2.0	
2	BMW	1 Series	2011	premium unleaded (required)	300.0	6.0	MANUAL	rear wheel drive	2.0	Li
3	BMW	1 Series	2011	premium unleaded (required)	230.0	6.0	MANUAL	rear wheel drive	2.0	
4	BMW	1 Series	2011	premium unleaded (required)	230.0	6.0	MANUAL	rear wheel drive	2.0	
...	
11909	Acura	ZDX	2012	premium unleaded (required)	300.0	6.0	AUTOMATIC	all wheel drive	4.0	Cros
11910	Acura	ZDX	2012	premium unleaded (required)	300.0	6.0	AUTOMATIC	all wheel drive	4.0	Cros
11911	Acura	ZDX	2012	premium unleaded (required)	300.0	6.0	AUTOMATIC	all wheel drive	4.0	Cros
11912	Acura	ZDX	2013	premium unleaded (recommended)	300.0	6.0	AUTOMATIC	all wheel drive	4.0	Cros
11913	Lincoln	Zephyr	2006	regular unleaded	221.0	6.0	AUTOMATIC	front wheel drive	4.0	

11914 rows × 16 columns

```
In [4]: train_raw_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11914 entries, 0 to 11913
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Make                   11914 non-null  object
1   Model                  11914 non-null  object
2   Year                   11914 non-null  int64
3   Engine Fuel Type      11911 non-null  object
4   Engine HP              11845 non-null  float64
5   Engine Cylinders      11884 non-null  float64
6   Transmission Type     11914 non-null  object
7   Driven_Wheels         11914 non-null  object
8   Number of Doors       11908 non-null  float64
9   Market Category       8172 non-null   object
10  Vehicle Size           11914 non-null  object
11  Vehicle Style          11914 non-null  object
12  highway MPG            11914 non-null  int64
13  city mpg               11914 non-null  int64
14  Popularity             11914 non-null  int64
15  MSRP                   11914 non-null  int64
dtypes: float64(3), int64(5), object(8)
memory usage: 1.5+ MB
```

```
In [5]: train_raw_df.dropna(inplace=True)
train_raw_df.reset_index(drop=True, inplace=True)
```

```
In [6]: train_raw_df["train_text"] = train_raw_df[['Market Category', 'Vehicle Size',
                                                'Vehicle Style']].apply(' '.join, axis=1)
```

```
In [7]: x_train = train_raw_df["train_text"]
y_train = train_raw_df.MSRP
```

```
In [8]: doc = " ".join(x_train)
```

```
In [9]: count_vec = CountVectorizer()
count_occurs = count_vec.fit_transform([doc])
count_occure_df = pd.DataFrame((count, word) for word, count in zip(count_occurs.data, count_vec.get_feature_names_out()))
count_occure_df.columns = ['Word', 'Count']
count_occure_df.sort_values('Count', ascending=False, inplace=True)
count_occure_df.head()
```

Out[9]:

	Word	Count
23	performance	3456
19	luxury	3279
20	midsize	3187
4	compact	3039
1	4dr	2771

Normalized Count Occurrence

If you think that high frequency may dominate the result and causing model bias. Normalization can be apply to pipeline easily.

```
In [10]: norm_count_vec = TfidfVectorizer(use_idf=False, norm='l2')
norm_count_occurs = norm_count_vec.fit_transform([doc])
norm_count_occur_df = pd.DataFrame((count, word) for word, count in zip(
    norm_count_occurs.toarray().tolist()[0], norm_count_vec.get_feature_names_out())
norm_count_occur_df.columns = ['Word', 'Count']
norm_count_occur_df.sort_values('Count', ascending=False, inplace=True)
norm_count_occur_df.head()
```

Out[10]:

	Word	Count
23	performance	0.386670
19	luxury	0.366867
20	midsize	0.356573
4	compact	0.340015
1	4dr	0.310030

TF-IDF

TF-IDF take another approach which is believe that high frequency may not able to provide much information gain. In another word, rare words contribute more weights to the model.

Word importance will be increased if the number of occurrence within same document (i.e. training record). On the other hand, it will be decreased if it occurs in corpus (i.e. other training records).

```
In [11]: tfidf_vec = TfidfVectorizer()
tfidf_count_occurs = tfidf_vec.fit_transform([doc])
tfidf_count_occur_df = pd.DataFrame((count, word) for word, count in zip(
    tfidf_count_occurs.toarray().tolist()[0], tfidf_vec.get_feature_names_out())
tfidf_count_occur_df.columns = ['Word', 'Count']
tfidf_count_occur_df.sort_values('Count', ascending=False, inplace=True)
tfidf_count_occur_df.head()
```

Out[11]:

	Word	Count
23	performance	0.386670
19	luxury	0.366867
20	midsize	0.356573
4	compact	0.340015
1	4dr	0.310030

Word2Vec

Word Embedding is a language modeling technique used for mapping words to vectors of real numbers. It represents words or phrases in vector space with several dimensions. Word embeddings can be generated using various methods like neural networks, co-occurrence matrix, probabilistic models, etc.

Word2Vec consists of models for generating word embedding. These models are shallow two-layer neural networks having one input layer, one hidden layer, and one output layer.

Word2Vec utilizes two architectures :

1. CBOW (Continuous Bag of Words): CBOW model predicts the current word given context words within a specific window. The input layer contains the context words and the output layer contains the current word. The hidden layer contains the number of dimensions in which we want to represent the current

word present at the output layer.

2. Skip Gram : Skip gram predicts the surrounding context words within specific window given current word. The input layer contains the current word and the output layer contains the context words. The hidden layer contains the number of dimensions in which we want to represent current word present at the input layer.

```
In [12]: data = []

# iterate through each sentence in the file
for i in x_train:
    temp = []

    # tokenize the sentence into words
    for j in word_tokenize(i):
        temp.append(j.lower())

    data.append(temp)
```

```
In [13]: # Create CBOW model
model1 = gensim.models.Word2Vec(data, min_count = 1, vector_size = 100,
                                window = 5)

# Print results
print("Cosine similarity between 'Luxury' and 'Performance' - CBOW : ",
      model1.wv.similarity('luxury', 'performance'))

print("Cosine similarity between 'Crossover' and 'Midsize' - CBOW : ",
      model1.wv.similarity('crossover', 'midsize'))
```

```
Cosine similarity between 'Luxury' and 'Performance' - CBOW : 0.95199424
Cosine similarity between 'Crossover' and 'Midsize' - CBOW : 0.9288742
```

```
In [14]: # Create Skip Gram model
model2 = gensim.models.Word2Vec(data, min_count = 1, vector_size = 100,
                                window = 5, sg = 1)

# Print results
print("Cosine similarity between 'Luxury' and 'Performance' - Skip Gram : ",
      model2.wv.similarity('luxury', 'performance'))

print("Cosine similarity between 'Crossover' and 'Midsize' - Skip Gram : ",
      model2.wv.similarity('crossover', 'midsize'))
```

```
Cosine similarity between 'Luxury' and 'Performance' - Skip Gram : 0.95101005
Cosine similarity between 'Crossover' and 'Midsize' - Skip Gram : 0.8570327
```